

# Building a Smarter AI powered spam classifier

## ***TEAM SPARROW***

### OVERVIEW:

In this project, we will be covering a simple approach to email classification(spam or not spam) using BERT

Steps are:

- We will load our data - mainly sentences and labels-span or not spam
- Load these in bert to generate an contextualized embedding vector of length 768
- - We will first apply preprocessing using the preprocessor object , refer the documentation
- - We will pass this preprocessed text to our model to generate the contextutailized embedding vector
- Finally pass this embedding vector to single neuron in output to do binary classificaton
- For maximizing performance we will be balancing our dataset and use a dropout layer to regularize the model and prevent overfitting

```
from google.colab import drive
drive.mount('/content/drive')
```

```
Drive already mounted at /content/drive; to attempt to forcibly
remount, call drive.mount("/content/drive", force_remount=True).
```

## Loading Dependencies

Includes

- Tensorflow\_hub : Place where all tenseorflow pretrained models are stored.
- Pandas : For data loading, manipulation and wrangling.
- Tensorflow\_text : Allows additional NLP text processing capabilities outside scope of tensorflow
- Skelarn : For doing data evaluation and splitting
- Matplotlib : For visualizaton

```
# installing tensorflow_text
```

```
!pip install tensorflow-text
```

```
Collecting tensorflow-text
```

```
  Downloading tensorflow_text-2.14.0-cp310-cp310-
```

```
manylinux_2_17_x86_64.manylinux2014_x86_64.whl (6.5 MB)
```

```
----- 6.5/6.5 MB 50.1 MB/s eta
```

```
0:00:00
```

```
Requirement already satisfied: tensorflow-hub>=0.13.0 in
```

```
/usr/local/lib/python3.10/dist-packages (from tensorflow-text)
```

```
(0.15.0)
```

```
Requirement already satisfied: tensorflow<2.15,>=2.14.0 in
```

```
/usr/local/lib/python3.10/dist-packages (from tensorflow-text)
```

```
(2.14.0)
```

```
Requirement already satisfied: absl-py>=1.0.0 in
```

```
/usr/local/lib/python3.10/dist-packages (from
```

```
tensorflow<2.15,>=2.14.0->tensorflow-text) (1.4.0)
```

```
Requirement already satisfied: astunparse>=1.6.0 in
```

```
/usr/local/lib/python3.10/dist-packages (from
```

```
tensorflow<2.15,>=2.14.0->tensorflow-text) (1.6.3)
```

```
Requirement already satisfied: flatbuffers>=23.5.26 in
```

```
/usr/local/lib/python3.10/dist-packages (from
```

```
tensorflow<2.15,>=2.14.0->tensorflow-text) (23.5.26)
```

```
Requirement already satisfied: gast!=0.5.0,!0.5.1,!0.5.2,>=0.2.1
```

```
in /usr/local/lib/python3.10/dist-packages (from
```

```
tensorflow<2.15,>=2.14.0->tensorflow-text) (0.5.4)
```

```
Requirement already satisfied: google-pasta>=0.1.1 in
```

```
/usr/local/lib/python3.10/dist-packages (from
```

```
tensorflow<2.15,>=2.14.0->tensorflow-text) (0.2.0)
```

```
Requirement already satisfied: h5py>=2.9.0 in
```

```
/usr/local/lib/python3.10/dist-packages (from
```

```
tensorflow<2.15,>=2.14.0->tensorflow-text) (3.9.0)
```

```
Requirement already satisfied: libclang>=13.0.0 in
```

```
/usr/local/lib/python3.10/dist-packages (from
```

```
tensorflow<2.15,>=2.14.0->tensorflow-text) (16.0.6)
```

```
Requirement already satisfied: ml-dtypes==0.2.0 in
```

```
/usr/local/lib/python3.10/dist-packages (from
```

```
tensorflow<2.15,>=2.14.0->tensorflow-text) (0.2.0)
```

```
Requirement already satisfied: numpy>=1.23.5 in
```

```
/usr/local/lib/python3.10/dist-packages (from
```

```
tensorflow<2.15,>=2.14.0->tensorflow-text) (1.23.5)
```

```
Requirement already satisfied: opt-einsum>=2.3.2 in
```

```
/usr/local/lib/python3.10/dist-packages (from
```

```
tensorflow<2.15,>=2.14.0->tensorflow-text) (3.3.0)
```

```
Requirement already satisfied: packaging in
```

```
/usr/local/lib/python3.10/dist-packages (from
```

```
tensorflow<2.15,>=2.14.0->tensorflow-text) (23.2)
```

```
Requirement already satisfied: protobuf!=4.21.0,!4.21.1,!4.21.2,!4.21.3,!4.21.4,!4.21.5,<5.0.0dev,>=3.20.3 in
```

```
/usr/local/lib/python3.10/dist-packages (from
```

tensorflow<2.15,>=2.14.0->tensorflow-text) (3.20.3)  
Requirement already satisfied: setuptools in  
/usr/local/lib/python3.10/dist-packages (from  
tensorflow<2.15,>=2.14.0->tensorflow-text) (67.7.2)  
Requirement already satisfied: six>=1.12.0 in  
/usr/local/lib/python3.10/dist-packages (from  
tensorflow<2.15,>=2.14.0->tensorflow-text) (1.16.0)  
Requirement already satisfied: termcolor>=1.1.0 in  
/usr/local/lib/python3.10/dist-packages (from  
tensorflow<2.15,>=2.14.0->tensorflow-text) (2.3.0)  
Requirement already satisfied: typing-extensions>=3.6.6 in  
/usr/local/lib/python3.10/dist-packages (from  
tensorflow<2.15,>=2.14.0->tensorflow-text) (4.5.0)  
Requirement already satisfied: wrapt<1.15,>=1.11.0 in  
/usr/local/lib/python3.10/dist-packages (from  
tensorflow<2.15,>=2.14.0->tensorflow-text) (1.14.1)  
Requirement already satisfied: tensorflow-io-gcs-filesystem>=0.23.1 in  
/usr/local/lib/python3.10/dist-packages (from  
tensorflow<2.15,>=2.14.0->tensorflow-text) (0.34.0)  
Requirement already satisfied: grpcio<2.0,>=1.24.3 in  
/usr/local/lib/python3.10/dist-packages (from  
tensorflow<2.15,>=2.14.0->tensorflow-text) (1.59.0)  
Requirement already satisfied: tensorboard<2.15,>=2.14 in  
/usr/local/lib/python3.10/dist-packages (from  
tensorflow<2.15,>=2.14.0->tensorflow-text) (2.14.1)  
Requirement already satisfied: tensorflow-estimator<2.15,>=2.14.0  
in /usr/local/lib/python3.10/dist-packages (from  
tensorflow<2.15,>=2.14.0->tensorflow-text) (2.14.0)  
Requirement already satisfied: keras<2.15,>=2.14.0 in  
/usr/local/lib/python3.10/dist-packages (from  
tensorflow<2.15,>=2.14.0->tensorflow-text) (2.14.0)  
Requirement already satisfied: wheel<1.0,>=0.23.0 in  
/usr/local/lib/python3.10/dist-packages (from astunparse>=1.6.0-  
>tensorflow<2.15,>=2.14.0->tensorflow-text) (0.41.2)  
Requirement already satisfied: google-auth<3,>=1.6.3 in  
/usr/local/lib/python3.10/dist-packages (from tensorboard<2.15,>=2.14-  
>tensorflow<2.15,>=2.14.0->tensorflow-text) (2.17.3)  
Requirement already satisfied: google-auth-oauthlib<1.1,>=0.5 in  
/usr/local/lib/python3.10/dist-packages (from tensorboard<2.15,>=2.14-  
>tensorflow<2.15,>=2.14.0->tensorflow-text) (1.0.0)  
Requirement already satisfied: markdown>=2.6.8 in  
/usr/local/lib/python3.10/dist-packages (from tensorboard<2.15,>=2.14-  
>tensorflow<2.15,>=2.14.0->tensorflow-text) (3.5)  
Requirement already satisfied: requests<3,>=2.21.0 in  
/usr/local/lib/python3.10/dist-packages (from tensorboard<2.15,>=2.14-  
>tensorflow<2.15,>=2.14.0->tensorflow-text) (2.31.0)  
Requirement already satisfied: tensorboard-data-server<0.8.0,>=0.7.0  
in /usr/local/lib/python3.10/dist-packages (from  
tensorboard<2.15,>=2.14->tensorflow<2.15,>=2.14.0->tensorflow-text)

(0.7.2)  
Requirement already satisfied: werkzeug>=1.0.1 in  
/usr/local/lib/python3.10/dist-packages (from tensorboard<2.15,>=2.14->tensorflow<2.15,>=2.14.0->tensorflow-text) (3.0.1)  
Requirement already satisfied: cachetools<6.0,>=2.0.0 in  
/usr/local/lib/python3.10/dist-packages (from google-auth<3,>=1.6.3->tensorboard<2.15,>=2.14->tensorflow<2.15,>=2.14.0->tensorflow-text) (5.3.2)  
Requirement already satisfied: pyasn1-modules>=0.2.1 in  
/usr/local/lib/python3.10/dist-packages (from google-auth<3,>=1.6.3->tensorboard<2.15,>=2.14->tensorflow<2.15,>=2.14.0->tensorflow-text) (0.3.0)  
Requirement already satisfied: rsa<5,>=3.1.4 in  
/usr/local/lib/python3.10/dist-packages (from google-auth<3,>=1.6.3->tensorboard<2.15,>=2.14->tensorflow<2.15,>=2.14.0->tensorflow-text) (4.9)  
Requirement already satisfied: requests-oauthlib>=0.7.0 in  
/usr/local/lib/python3.10/dist-packages (from google-auth-oauthlib<1.1,>=0.5->tensorboard<2.15,>=2.14->tensorflow<2.15,>=2.14.0->tensorflow-text) (1.3.1)  
Requirement already satisfied: charset-normalizer<4,>=2 in  
/usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->tensorboard<2.15,>=2.14->tensorflow<2.15,>=2.14.0->tensorflow-text) (3.3.1)  
Requirement already satisfied: idna<4,>=2.5 in  
/usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->tensorboard<2.15,>=2.14->tensorflow<2.15,>=2.14.0->tensorflow-text) (3.4)  
Requirement already satisfied: urllib3<3,>=1.21.1 in  
/usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->tensorboard<2.15,>=2.14->tensorflow<2.15,>=2.14.0->tensorflow-text) (2.0.7)  
Requirement already satisfied: certifi>=2017.4.17 in  
/usr/local/lib/python3.10/dist-packages (from requests<3,>=2.21.0->tensorboard<2.15,>=2.14->tensorflow<2.15,>=2.14.0->tensorflow-text) (2023.7.22)  
Requirement already satisfied: MarkupSafe>=2.1.1 in  
/usr/local/lib/python3.10/dist-packages (from werkzeug>=1.0.1->tensorboard<2.15,>=2.14->tensorflow<2.15,>=2.14.0->tensorflow-text) (2.1.3)  
Requirement already satisfied: pyasn1<0.6.0,>=0.4.6 in  
/usr/local/lib/python3.10/dist-packages (from pyasn1-modules>=0.2.1->google-auth<3,>=1.6.3->tensorboard<2.15,>=2.14->tensorflow<2.15,>=2.14.0->tensorflow-text) (0.5.0)  
Requirement already satisfied: oauthlib>=3.0.0 in  
/usr/local/lib/python3.10/dist-packages (from requests-oauthlib>=0.7.0->google-auth-oauthlib<1.1,>=0.5->tensorboard<2.15,>=2.14->tensorflow<2.15,>=2.14.0->tensorflow-text) (3.2.2)

```
Installing collected packages: tensorflow-text
Successfully installed tensorflow-text-2.14.0
```

```
import tensorflow_hub as hub
import pandas as pd
import tensorflow_text as text
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
import numpy as np
```

## Loading Data

- Read Data
- Display data USING PANDAS

```
# load data
df = pd.read_csv('/content/drive/MyDrive/CLG_VK_18_/spam_data.csv')
df.head()
```

|   | Category | Message   |
|---|----------|---|
| 0 | ham      | Go until jurong point, crazy.. Available only ... |
| 1 | ham      | Ok lar... Joking wif u oni...                     |
| 2 | spam     | Free entry in 2 a wkly comp to win FA Cup fina... |
| 3 | ham      | U dun say so early hor... U c already then say... |
| 4 | ham      | Nah I don't think he goes to usf, he lives aro... |

## Data Analysis

- Check the description by grouping by category :
- no of data points for each category - count
- no of unique values in each category - unique

```
# check count and unique and top values and their frequency
df['Category'].value_counts()

ham      4825
spam      747
Name: Category, dtype: int64
```

Clearly dataset is imbalanced - not so much but still it can affect our model. Need to use some type of regularization like downsampling dataset for majority class

## Downsampling Dataset

Includes:

- Check percentage of unbalances.
- Creating 2 new dataframes out of existing one.
- Taking any random minority no of samples - (747) for majority class (4825).
- Creating a balanced dataset by concating 2 new data frames.

```
# check percentange of data - states how much data needs to be
balanced
str(round(747/4825,2))+ '%'

{"type": "string"}

# creating 2 new dataframe as df_ham , df_spam

df_spam = df[df['Category']=='spam']
print("Spam Dataset Shape:", df_spam.shape)

df_ham = df[df['Category']=='ham']
print("Ham Dataset Shape:", df_ham.shape)

Spam Dataset Shape: (747, 2)
Ham Dataset Shape: (4825, 2)

# downsampling ham dataset - take only random 747 example
# will use df_spam.shape[0] - 747

df_ham_downsampled = df_ham.sample(df_spam.shape[0])
df_ham_downsampled.shape

(747, 2)

# concating both dataset - df_spam and df_ham_balanced to create
df_balanced dataset
df_balanced = pd.concat([df_spam , df_ham_downsampled])
df_balanced.head()
```

|    | Category | Message   |
|----|----------|---|
| 2  | spam     | Free entry in 2 a wkly comp to win FA Cup fina... |
| 5  | spam     | FreeMsg Hey there darling it's been 3 week's n... |
| 8  | spam     | WINNER!! As a valued network customer you have... |
| 9  | spam     | Had your mobile 11 months or more? U R entitle... |
| 11 | spam     | SIX chances to win CASH! From 100 to 20,000 po... |

```
df_balanced['Category'].value_counts()

spam      747
ham       747
Name: Category, dtype: int64

df_balanced.sample(10)
```

|      | Category | Message   |
|------|----------|---|
| 5041 | spam     | Natalie (20/F) is inviting you to be her frien... |

|      |      |   |
|------|------|---|
| 3713 | ham  | Wat u doing there?                                |
| 1172 | spam | Got what it takes 2 take part in the WRC Rally... |
| 1217 | spam | You have 1 new voicemail. Please call 08719181... |
| 403  | ham  | The hair cream has not been shipped.              |
| 4863 | spam | **FREE MESSAGE**Thanks for using the Auction S... |
| 2503 | ham  | Ola would get back to you maybe not today but ... |
| 5267 | ham  | Anything lar then ü not going home 4 dinner?      |
| 2719 | spam | 18 days to Euro2004 kickoff! U will be kept in... |
| 3678 | ham  | Great! So what attracts you to the brothas?       |

## Data Prepration

1. Create Numerical Repersentation Of Category - One hot encoding
  - Create a new column
  - Use `df[col].apply(lambda function)`
  - Lambda Function - if spam return 1, else return 0 (for ham) - ternary operators : `[lambda x : value expression else value]`

```
# creating numerical repersentation of category - one hot encoding
df_balanced['spam'] = df_balanced['Category'].apply(lambda x:1 if
x=='spam' else 0)
```

```
# displaying data - spam -1 , ham-0
df_balanced.sample(4)
```

|      | Category | Message   | spam |
|------|----------|---|------|
| 928  | ham      | K:)i will give my kvb acc details:)               | 0    |
| 435  | ham      | You available now? I'm like right around hills... | 0    |
| 2760 | ham      | Can meh? Thgt some will clash... Really ah, i ... | 0    |
| 4682 | ham      | Are you staying in town ?                         | 0    |

1. Do train-test split
  - split dataset into 80-20 ratio with 80% train and remaing as test
  - for eveness of data we will use `stratify` agrument which ensures same ratio of both category is loaded for each case, even if one catagoy has more training samples - prevents overfitting

Store our data in:

- `X_train, y_train` - training set(training\_data and labels respectively)
- `X_test,, y_test` - testing set(testing\_data and labels)

```
# loading train test split
from sklearn.model_selection import train_test_split

X_train, X_test , y_train, y_test =
train_test_split(df_balanced['Message'], df_balanced['spam'],
```

```

df_balanced['spam'])
# check for stratification
y_train.value_counts()
1      560
0      560
Name: spam, dtype: int64
560/560
1.0
y_test.value_counts()
1      187
0      187
Name: spam, dtype: int64
187/187
1.0
stratify =

```

*-> Almost similar, means data is downsampled now*

## Model Creation

Our Model is BERT , which will do 2 thing:

- Preprocess our training data that will be feeded - includes **adding additional token CLF, PAD and SEP** to generate `input_mask`, `input_type_ids`, `input_word_ids` (token given to each word in sentences)
- Note: no of words in sentence - 128/ max length of sentence can be 128

## Downloading BERT

Model specification :

- Layers - 12
  - Hidden layers - 768 - embedding size
  - Attention - 12 Name - Bert Small --- This model has 2 parts:
  - Bert\_preprocessor - preprocess the text to be BERT ready
  - Bert\_encoder - do the actual encoding
- Steps:



Preprocessor

- create a keras hub layer from the preprocessing url

Encoder

- create a keras hub layer from the encoder/ model url

Awesome functionality provided by Tf hub API

+

Creating our own model using functional model api- link old layers to new layers rather than building it(in a sequential way) and allows sharing of layers too

Info:

- Text the embedding as input - text\_input
- Create a Single output dense layer
- Add dropout to reduce overfitting

```
# downloading preprocessing files and model
bert_preprocessor =
hub.KerasLayer('https://tfhub.dev/tensorflow/bert_en_uncased_preprocess/3')
bert_encoder =
hub.KerasLayer('https://tfhub.dev/tensorflow/bert_en_uncased_L-12_H-768_A-12/4')
```

## Process And Encode Data

Use functional API to process and encode data in the layers itself

- Create a input layers with shape(), type - tf.string, and layer name - text - `TEXT_INPUT`
- Pass `TEXT_INPUT` into bert\_preprocessor - `PREPROCESSED TEXT[*]`
- Pass the above[\*] to encoder - `EMBEED`
- pass pooled\_outputs of EMBEED to dropout layer - `DROPOUT`
- create a dense layer with activation as `sigmoid OUTPUTS`
- Create out MODEL (inputs - text\_input, outputs - dropout)

```
import tensorflow as tf

text_input = tf.keras.layers.Input(shape = (), dtype = tf.string, name = 'Inputs')
preprocessed_text = bert_preprocessor(text_input)
embed = bert_encoder(preprocessed_text)
```

```
dropout = tf.keras.layers.Dropout(0.1, name = 'Dropout')
(embded['pooled_output'])
outputs = tf.keras.layers.Dense(1, activation = 'sigmoid', name =
'Dense')(dropout)
```

```
# creating final model
```

```
model = tf.keras.Model(inputs = [text_input], outputs = [outputs])
```

```
# check summary of model
```

```
model.summary()
```

```
Model: "model_2"
```

| Layer (type)<br>Connected to   | Output Shape   | Param #       |
|--|--|---------------|
| =====  |  |               |
| Inputs (InputLayer)  | [(None,)]  | 0 []          |
| <br>   |  |               |
| keras_layer_2 (KerasLayer) ['Inputs[0][0]']  | {'input_type_ids': (None, 128),<br><br>'input_mask': (None, 128),<br><br>, 'input_word_ids': (None, 128)}  | 0             |
| <br>   |  |               |
| keras_layer_3 (KerasLayer) ['keras_layer_2[0][0]',<br>'keras_layer_2[0][1]',<br>'keras_layer_2[0][2]'] | {'encoder_outputs': [(None, 128, 768),<br>(None, 128, 768),<br>(None, 128, 768),<br>(None, 128, 768),<br>(None, 128, 768),<br>(None, 128, 768),<br>(None, 128, 768),<br>(None, 128, 768)]} | 1094822<br>41 |

```

        (None, 128, 768),
        (None, 128, 768),
        (None, 128, 768),
        (None, 128, 768)],
        'default': (None, 768),
        'sequence_output': (None,
        128, 768),
        'pooled_output': (None, 7
        68)}

Dropout (Dropout)          (None, 768)          0
['keras_layer_3[0][13]']

Dense (Dense)              (None, 1)          769
['Dropout[0][0]']

=====
=====
Total params: 109483010 (417.64 MB)
Trainable params: 769 (3.00 KB)
Non-trainable params: 109482241 (417.64 MB)

```

## Compiling model

- Optimizer - ADAM
- Loss - binary\_crossentropy
- metrics - accuracy, precision and recall

```

Metrics = [tf.keras.metrics.BinaryAccuracy(name = 'accuracy'),
            tf.keras.metrics.Precision(name = 'precision'),
            tf.keras.metrics.Recall(name = 'recall')
            ]

model.compile(optimizer='adam',
              loss = 'binary_crossentropy',
              metrics = Metrics)

```

```
#@title Optional
# optional - defining tensorflow callbacks
import tensorflow as tf
import datetime
%load_ext tensorboard

!rm -rf ./logs/
log_dir = "logs/fit/" + datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
tensorboard_callback=tf.keras.callbacks.TensorBoard(log_dir=log_dir,
histogram_freq=1)

The tensorboard extension is already loaded. To reload it, use:
%reload_ext tensorboard
```

## Training Model

- Recommended to use GPU - providing so many training data
- We training our model on training set
- For 10 epochs only - so model don't overfit - given enough training data

```
%tensorboard --logdir logs/fit

Reusing TensorBoard on port 6006 (pid 3229), started 1:03:18 ago. (Use
'!kill 3229' to kill it.)

<IPython.core.display.Javascript object>

history = model.fit(X_train, y_train, epochs = 10 , callbacks =
[tensorboard_callback])

Epoch 1/10
35/35 [=====] - 48s 1s/step - loss: 0.6358 -
accuracy: 0.6500 - precision: 0.6567 - recall: 0.6286
Epoch 2/10
35/35 [=====] - 38s 1s/step - loss: 0.4922 -
accuracy: 0.8286 - precision: 0.8345 - recall: 0.8196
Epoch 3/10
35/35 [=====] - 37s 1s/step - loss: 0.4205 -
accuracy: 0.8607 - precision: 0.8435 - recall: 0.8857
Epoch 4/10
35/35 [=====] - 38s 1s/step - loss: 0.3736 -
accuracy: 0.8768 - precision: 0.8728 - recall: 0.8821
Epoch 5/10
35/35 [=====] - 38s 1s/step - loss: 0.3400 -
accuracy: 0.8955 - precision: 0.8866 - recall: 0.9071
Epoch 6/10
35/35 [=====] - 37s 1s/step - loss: 0.3129 -
accuracy: 0.8991 - precision: 0.8914 - recall: 0.9089
```

```

Epoch 7/10
35/35 [=====] - 38s 1s/step - loss: 0.2966 -
accuracy: 0.9045 - precision: 0.9009 - recall: 0.9089
Epoch 8/10
35/35 [=====] - 39s 1s/step - loss: 0.2790 -
accuracy: 0.9143 - precision: 0.9042 - recall: 0.9268
Epoch 9/10
35/35 [=====] - 37s 1s/step - loss: 0.2674 -
accuracy: 0.9134 - precision: 0.9069 - recall: 0.9214
Epoch 10/10
35/35 [=====] - 38s 1s/step - loss: 0.2560 -
accuracy: 0.9116 - precision: 0.9109 - recall: 0.9125

```

## Model Evaluation

- Evaluating model performance using `model.evaluate(X_test, y_test)`
- Predicting `X_test` - `y_pred` -- Checking its values as 1 or 0
- Getting Confusion matrix -- Flattening `y_pred` -- Plotting confusion matrix
- Getting classification report

```

# Evaluating performace
model.evaluate(X_test,y_test)

12/12 [=====] - 5s 359ms/step - loss: 0.2729
- accuracy: 0.9198 - precision: 0.9026 - recall: 0.9412

[0.272863507270813, 0.9197860956192017, 0.9025641083717346,
0.9411764740943909]

# getting y_pred by predicting over X_text and flattening it
y_pred = model.predict(X_test)
y_pred = y_pred.flatten() # require to be in one dimensional array ,
for easy manipulation

12/12 [=====] - 4s 335ms/step

# checking the results y_pred
import numpy as np

y_pred = np.where(y_pred>0.5,1,0 )
y_pred

array([1, 1, 1, 0, 1, 0, 1, 1, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 1, 0, 0,
0,
      0, 0, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 1, 0, 1, 0, 1, 1, 1, 1,
1,

```

```

1,      1, 1, 1, 1, 0, 0, 0, 0, 1, 1, 0, 0, 1, 1, 1, 0, 1, 0, 1, 1, 1,
1,      0, 0, 1, 1, 1, 0, 1, 0, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0, 1, 1, 0, 0,
1,      0, 0, 1, 1, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 0, 1, 1, 1, 0, 0, 1,
1,      1, 0, 0, 0, 0, 1, 0, 0, 0, 1, 1, 0, 0, 0, 1, 0, 1, 0, 0, 1, 1,
1,      0, 1, 0, 1, 0, 0, 0, 1, 0, 1, 0, 1, 1, 0, 1, 0, 0, 0, 0, 1, 0,
0,      1, 1, 1, 1, 1, 1, 1, 1, 0, 0, 1, 0, 0, 0, 1, 0, 1, 1, 0, 0, 0,
1,      1, 0, 1, 1, 0, 1, 1, 1, 1, 1, 1, 0, 0, 1, 1, 0, 0, 1, 0, 1, 1,
1,      0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 0, 1, 0, 0, 0, 0, 0, 1, 1, 1, 1,
0,      0, 0, 1, 1, 0, 1, 1, 1, 0, 1, 0, 1, 1, 0, 0, 1, 1, 1, 0, 0, 0,
1,      1, 0, 1, 1, 0, 1, 0, 1, 1, 1, 0, 1, 1, 1, 1, 0, 1, 1, 0, 0, 0,
1,      0, 1, 1, 0, 0, 1, 0, 1, 0, 1, 0, 0, 0, 1, 1, 1, 0, 1, 1, 1, 0,
1,      1, 1, 1, 0, 0, 0, 1, 1, 1, 1, 0, 1, 0, 0, 1, 0, 0, 1, 0, 1, 0,
0,      1, 1, 0, 0, 1, 1, 1, 1, 0, 0, 1, 0, 1, 0, 1, 1, 1, 1, 1, 0, 1,
0,      1, 0, 0, 1, 1, 0, 0, 1, 0, 0, 1, 0, 0, 0, 1, 1, 1, 0, 1, 0, 0,
0,      0, 1, 0, 0, 0, 1, 0, 0, 1, 0, 1, 1, 1, 0, 1, 0, 0, 0, 1, 1, 0,
1])

```

**Not so understandable so plotting confusion matrix and classification report for good visualization**

```

# importing confusion maxtrix
from sklearn.metrics import confusion_matrix , classification_report

# creating confusion matrix
cm = confusion_matrix(y_test,y_pred)
cm


array([[168,  19],
       [ 11, 176]])

# plotting as graph - importing seaborn
import seaborn as sns

# creating a graph out of confusion matrix
sns.heatmap(cm, annot = True, fmt = 'd')

```

```
plt.xlabel('Predicted')
plt.ylabel('Actual')
Text(50.72222222222214, 0.5, 'Actual')
```



```
# printing classification report
print(classification_report(y_test , y_pred))
```

|              | precision | recall | f1-score | support |
|--------------|-----------|--------|----------|---------|
| 0            | 0.94      | 0.90   | 0.92     | 187     |
| 1            | 0.90      | 0.94   | 0.92     | 187     |
| accuracy     |           |        | 0.92     | 374     |
| macro avg    | 0.92      | 0.92   | 0.92     | 374     |
| weighted avg | 0.92      | 0.92   | 0.92     | 374     |

Good Precesion And Recall Score, but can be improved

## Model Prediction

- We will be predicting data on text coprus, value > 5 is most likely be spam

```

predict_text = [
    # Spam
    'URGENT! You have won a 1 week FREE membership in our
    £100,000 Prize Jackpot! Txt the word: CLAIM to No: 81010 T&C
    www.dbuk.net LCCLTD POBOX 4403LDNW1A7RW18',
    'WINNER!! As a valued network customer you have been
    selected to receivea £900 prize reward! To claim call 09061701461.
    Claim code KL341. Valid 12 hours only.',
    'England v Macedonia - dont miss the goals/team news.
    Txt ur national team to 87077 eg ENGLAND to 87077 Try:WALES, SCOTLAND
    4txt/ú1.20 POBOXox36504W45WQ 16+',

    #ham
    'U still going to the mall?',
    'Haha awesome, be there in a minute.',
    'Shit that is really shocking and scary, cant imagine
    for a second. Def up for night out. Do u think there is somewhere i
    could crash for night, save on taxi?'
]

test_results = model.predict(predict_text)

1/1 [=====] - 1s 799ms/step

output = np.where(test_results>0.5, 'spam', 'ham')

output
array([[ 'spam'],
       [ 'spam'],
       [ 'spam'],
       [ 'ham'],
       [ 'ham'],
       [ 'ham']], dtype='<U4')

```

## Additional Content

- Create a function which will take in `sentece array` and return the embedding vector for entire sentece - `pooled_output`  
 STEPS: To do so inside the we follow 3 steps:
  1. We pass the `sentence array` to `bert_preprocessor` as it can act a function point and name it **preprocessed\_text**
  2. Now we pass this `preprocessed sentence` into `encoder` and it return a embedding vector dictionary



3. We return only the `pooled_output` as we are interested in only the entire sentence encoding
- 

Later we compare the embedding vector using `cosine - similarity` from `sklearn.metrics.pairwise` class

```
def get_embedding(sentence_arr):  
    'takes in sentence array and return embedding vector'  
    preprocessed_text = bert_preprocessor(sentence_arr)  
    embeddings = bert_encoder(preprocessed_text)['pooled_output']  
    return embeddings  
  
e = get_embedding([  
    'We'd all like to get a $10,000 deposit on our bank  
    accounts out of the blue, but winning a prize—especially if you've  
    never entered a contest',  
    'The image you sent is a UI bug, I can check that your  
    article is marked as regular and is not in the monetization program.'  
])  
  
# load similartiy score  
from sklearn.metrics.pairwise import cosine_similarity  
  
# check similarity score  
print(f'Similarity score between 1st sentence(spam) and second  
sentence(spam) : {cosine_similarity([e[0]] , [e[1]])}')  
  
Similarity score between 1st sentence(spam) and second  
sentence(spam) : [[0.853919]]
```

- Not exact similarity, may show unexpected results as can be seen - they are somewhat similar but it's false as spam and actual can't be same